

Article Presentation

Mathematics of Deep Learning

Authors: René Vidal, Joan Bruna, Raja Giryes and Stefano Soatto

Juan Carrascal, Maria Perpiñán, Edward Soto, Mayra Vega

Universidad Nacional de Colombia

May 2023

- 1 Introduction
- 2 Global Optimality in Deep Learning
- 3 Geometric Stability in Deep Learning
- 4 Structure Based Theory for Deep Learning
- 5 Towards an Information-Theoretic Framework
- 6 Bibliography

The paper reviews recent work that aims to provide a mathematical justification for several properties of deep networks, such as global optimality, geometric stability, and invariance of the learned representations.

The structure of the paper is:

- The problem of training deep networks and conditions for global optimality.
- The invariance and stability properties of CNN.
- Structural properties of deep networks.
- Information-theoretic properties of deep representations.

Global Optimality in Deep Learning

We now study the problem of learning the parameters $W = \{W^l\}_{l=1}^L$ of a deep network from N training examples (X, Y) . The problem of learning the network weights W is formulated as the following optimization problem

$$\min_{\{W^l\}_{l=1}^L} \mathcal{L}(Y, \Phi(X, W^1, \dots, W^L)) + \lambda \Theta(W^1, \dots, W^L), \quad (1)$$

where $\mathcal{L}(Y, \Phi)$ is the loss function and Θ is the regularization term.

Global Optimality in Deep Learning

A. The challenge of non-convexity in neural network training

Key challenge: For most deep networks, $\mathcal{L}(Y, \Phi)$ is a convex function, however the map $\Phi(X, Y)$ is not due to the product of the W^l variables and the nonlinearities in of the activation functions.

Why is this a problem?




Fig 1: Example critical points of a non-convex function (shown in red). (a,c) Plateaus. (b,d) Global minima. (e,g) Local maxima. (f,h) Local minima.

Global Optimality in Deep Learning

A. The challenge of non-convexity in neural network training

Theoretical findings

In [1] they prove the following statements:

- For large-size networks, most local minima are equivalent and yield similar performance on a test set.
- The probability of finding a bad (high value) local minimum is non-zero for small-size networks and decreases quickly with network size.

Global Optimality in Deep Learning

B. Global optimality for positively homogeneous networks

Key challenge: Can we say something about the network optimality conditions without making any assumption on the input data distribution, the weight parameters or the network initialization?

Theoretical findings




Fig 2: Critical points distribution [2, 3].

Global Optimality in Deep Learning

B. Global optimality for positively homogeneous networks

Theoretical findings

Theorem

If Φ and Θ are the sum of positively homogeneous functions of the same degree, then any local minimizer of the non-convex optimization problem

$$\min_{\{W^l\}_{l=1}^L} \mathcal{L}(Y, \Phi_r(X, W^1, \dots, W^L)) + \lambda \sum_{i=1}^r \Theta(W^1, \dots, W^L), \quad (2)$$

such that $(W_{i_0}^1, \dots, W_{i_0}^L) = (0, \dots, 0)$ for some $i_0 \in \{1, \dots, r\}$ is a global minimizer of (2). Moreover, $X = \Phi_r(W^1, \dots, W^L)$ is a global minimizer of (1).

Geometric Stability in Deep Learning

Convolutional architectures are the key to the success of most deep learning vision models.

Fig 3: Convolutional filter

In these architectures, there is a notion of geometric stability which provides a possible framework to understand its success.

Geometric Stability in Deep Learning

Representation:

In image analysis applications, images can be thought of as functions on the unit square $\Omega = [0, 1]^2$. This representation is given by: $X \in L^2(\Omega)$

$$X : \Omega \longrightarrow \mathbb{R} \quad \text{with} \quad \int_{-\infty}^{\infty} |X(t)|^2 dt < \infty.$$

Let $f : L^2(\Omega) \longrightarrow \mathcal{Y}$ the unknown function we want to learn.

Stationarity:

Given a translation operator

$$\mathcal{T}_v X(u) = X(u - v) \quad u, v \in \Omega, \tag{3}$$

we may consider following properties:

Geometric Stability in Deep Learning

Invariance:

$f(\mathcal{T}_v X) = f(X)$ for any $X \in L^2(\Omega)$ and $v \in \Omega$. This is typically the case in object classification tasks.




Fig 4: Invariance principle

Geometric Stability in Deep Learning

Equivariance:

$f(\mathcal{T}_v X) = \mathcal{T}_v f(X)$ for any $X \in L^2(\Omega)$ and $v \in \Omega$. Observed in tasks of object localization and semantic segmentation.




Fig 5: Equivariance principle

Geometric Stability in Deep Learning

Local deformations:

Mathematically, given a smooth vector field $\tau : \Omega \longrightarrow \Omega$, if a function L_τ is such that:

$$L_\tau X(u) := X(u - \tau(u)), \quad (4)$$

we say that L_τ is a deformation. This kind of functions can model local translations, changes in viewpoint and rotations.




Fig 6: Local deformation

Geometric Stability in Deep Learning

Important facts:

- Most tasks studied in computer vision are not only translation invariant/equivariant, but, more importantly, also stable with respect to local deformations [4].
- Whereas long-range dependencies indeed exist in natural images and are critical to object recognition, they can be captured and down-sampled at different scales.
- CNNs strike a good balance in terms of approximation power, optimization, and invariance [5].
- Recently there has been an effort to extend the geometric stability priors to data that is not defined over an Euclidean domain [6].

Geometric Stability in Deep Learning

Example:

Brain tumor segmentation task makes use of an architecture called “*encoder-decoder*”. We can observe the mentioned geometric properties in this experiment.




Fig 7: Brain tumor segmentation task

Structure Based Theory for Deep Learning

Relationship between the structure of the data and generalization error

Key challenge: Understanding the good generalization observed in practice for deep networks with a large number of parameters or deep architectures.

The generalization error is then given as:

$$\text{GE}(\Phi) = |\ell_{\text{exp}}(\Phi) - \ell_{\text{emp}}(\Phi)|.$$

Where:

$$\ell_{\text{emp}}(\Phi) = \frac{1}{N} \sum_{X_i \in \mathcal{T}_N} \ell(Y_i, \Phi(X_i, W)),$$

$$\ell_{\text{exp}}(\Phi) = \mathbb{E}_{(X,Y) \sim P}[\ell(Y, \Phi(X, W))].$$

Structure Based Theory for Deep Learning

Relationship between the structure of the data and generalization error

Approaches: Measures such as **VC-dimension**, Rademacher or Gaussian complexities, and algorithm robustness have been used to bound the generalization error in deep networks.

Problem: Don't fully explain the good generalization observed in practice for DN with a large number of parameters or deep architectures.

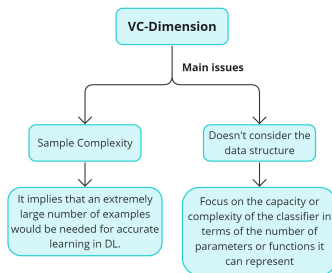


Fig 8: Issues with VC-dimension

Structure Based Theory for Deep Learning

Relationship between the structure of the data and generalization error

Solution: New approaches such as **Classification margin**.

Definition

The classification margin of a training sample $s_i = (\mathbf{x}_i, y_i)$ measured by a metric d is defined as

$$\gamma^d(s_i) = \sup \{a : d(\mathbf{x}_i, \mathbf{x}) \leq a \implies g(\mathbf{x}) = y_i \forall \mathbf{x}\}.$$




Fig 9: Sketch of the distortion of two classes with distinguishable angle

Structure Based Theory for Deep Learning

Relationship between the structure of the data and generalization error

Solution: New approaches such as **Classification margin**.

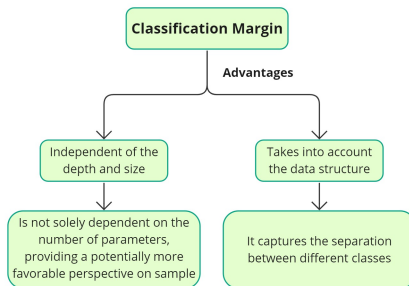


Fig 10: Advantages of classification margin

Towards an Information-Theoretic Framework

In multiclass classification problems the loss function in neural networks is the *empirical cross-entropy*.

Definition

The *empirical cross-entropy*, as a loss function, is defined as:

$$\ell(W) = \mathbb{E}_{P(X,Y)}(-\log(\Phi(X, W)))$$

It is well known that to solve overfitting or high variance problems, techniques like regularization are implemented. The authors propose two different loss functions to address these problems.

Towards an Information-Theoretic Framework

From an information theory point of view, a type of regularization can be restrict the stored information in the training weights [7].

Train regularization

Suppose $P(W)$ is a prior on the weights W . Define a loss by:

$$\ell(W) = H(Y|X, W) + \lambda KL(P(W|X, Y) || P(W))$$

Some problems were occurring in the attempt to calculate the Kullback-Leibler divergence. Recent advancements in the field of optimization have helped solve this problem.

Towards an Information-Theoretic Framework

It is well known that neural networks can adjust for noise in data. This idea can be explained as restricting the information extracted from the compressed representation of X by the network.

Test regularization

Suppose Z is a stochastic representation of X learned by a layer. Define a loss by:

$$\ell(W) = H(Y|Z, W) + \lambda I(Z; X)$$

Future work

Formally both approaches have no relation. But authors conjecture a relation exist and can be the explanation of generalization of networks. Bounds have been not found yet.

References I

- [1] Anna Choromanska, Mikael Hena, Michael Mathieu, Gerard Ben Arous, and Yann LeCun.
The loss surfaces of multilayer networks.
International Conference on Artificial Intelligence and Statistics, pages 192,204, 2015.
- [2] Benjamin D. Haeffele and René Vidal.
Global optimality in neural network training.
IEEE Conference on Computer Vision and Pattern Recognition, 2017.
- [3] Benjamin D. Haeffele and René Vidal.
Global optimality in tensor factorization, deep learning, and beyond.
arXiv, 2015.
- [4] Stéphane Mallat.
Understanding deep convolutional networks.
Phil. Trans. R. Soc. A, 2016.

References II

- [5] D. Ramanan P. Felzenszwalb R. Girshick, D. McAllester.
Object detection with discriminatively trained part-based models.
IEEE Transactions on Pattern Analysis and Machine Intelligence,
2010.
- [6] M. Bronstein, Y. LeCun J. Bruna, A. Szlam, and P. Vandergheynst.
Geometric deep learning: going beyond euclidean data.
arXiv preprint arXiv:1611.08097, 2016.
- [7] Geoffrey E Hinton and Drew Van Camp.
Keeping the neural networks simple by minimizing the description
length of the weights.
*In Proceedings of the sixth annual conference on Computational
learning theory*, pages 5–13, 1993.
- [8] Marco Gori and Alberto Tesi.
On the problem of local minima in backpropagation.
IEEE Transactions on Pattern Analysis and Machine Intelligence,
14(1):76,86, 1992.

References III

- [9] Yoshua Bengio, Nicolas L. Roux, Pascal Vincent, Olivier Delalleau, and Patrice Marcotte.
Convex neural networks.
Neural Information Processing Systems, pages 123,130, 2005.
- [10] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton.
Imagenet classification with deep convolutional neural networks.
Neural Information Processing Systems, pages 1097, 1105, 2012.
- [11] Jure Sokolić, Raja Giryes, Guillermo Sapiro, and Miguel R. D. Rodrigues.
Generalization error of deep neural networks: Role of classification margin and data structure.
In 2017 International Conference on Sampling Theory and Applications (SampTA), pages 147–151, 2017.

- [12] Raja Giryes, Guillermo Sapiro, and Alex M. Bronstein.
Deep neural networks with random gaussian weights: A universal
classification strategy?
IEEE Transactions on Signal Processing, 64(13):3444–3457, 2016.